


NIO Concepts

- **Channels**
Provide for data transfers between sockets and buffers
- **Buffers**
Contiguous extent of memory for processing data
- **Selector**
Allows to wait on several channels until one or more become available for data transfer



Fabienne Boyer, Basics of Distributed Programming

NIO – Server programming

```
public class NioServer implements Runnable {
    private InetAddress address;
    private int port, length = ...;
    private ServerSocketChannel serverChannel;
    private Selector selector;

    public NioServer(InetAddress hostAddress, int port) throws IOException {
        this.hostAddress = hostAddress;
        this.port = port;
        selector = SelectorProvider.provider().openSelector();

        // Create a new non-blocking server socket
        serverChannel = ServerSocketChannel.open();
        serverChannel.configureBlocking(false);
        serverChannel.socket().bind(new InetSocketAddress(hostAddress, port));

        // Be notified when connection requests arrive
        serverChannel.register(selector, SelectionKey.OP_ACCEPT);
    }

    public static void main(String[] args) {
        try { new Thread(new NioServer(null, 8888)).start(); } catch (IOException) ..
    }
}
```

2

NIO – Server programming

```
public void run() {
    while (true) {
        try {
            // Wait for an event one of the registered channels
            this.selector.select();

            // Some events have been received
            Iterator selectedKeys = this.selector.selectedKeys().iterator();
            while (selectedKeys.hasNext()) {
                SelectionKey key = (SelectionKey) selectedKeys.next();
                selectedKeys.remove();
                if (!key.isValid()) { continue; }

                // Handle the event
                if (key.isAcceptable()) { handleAccept(key); }
                } else if (key.isConnectable()) { handleConnect(key); }
                } else if (key.isReadable()) { handleRead(key); }
                } else if (key.isWritable()) { handleWrite(key); }
            } } catch (Exception e) { ... }
    }
}
```

Fabienne Boyer, Basics of Distributed Programming

3

NIO – Client programming

```
public class NioClient implements Runnable {
    private InetAddress address;
    private int port, length = ...;
    private SocketChannel clientChannel;
    private Selector selector;

    public NioClient(InetAddress hostAddress, int port) throws IOException {
        this.hostAddress = hostAddress;
        this.port = port;
        selector = SelectorProvider.provider().openSelector();

        // Create a new non-blocking socket channel
        clientChannel = SocketChannel.open();
        clientChannel.configureBlocking(false);

        // Be notified when connection is accepted
        clientChannel.register(selector, SelectionKey.OP_CONNECT);

        // Connect to the server
        clientChannel.connect(new InetSocketAddress(hostAddress, port));
    }

    public static void main(String[] args) {
        try { new Thread(new NioClient(null, 8888)).start(); } catch (IOException) ..
    }
}
```

4

NIO – Client programming

```
public void run() {
    while (true) {
        try {
            // Wait for an event one of the registered channels
            this.selector.select();

            // Some events have been received
            Iterator selectedKeys = this.selector.selectedKeys().iterator();
            while (selectedKeys.hasNext()) {
                SelectionKey key = (SelectionKey) selectedKeys.next();
                selectedKeys.remove();
                if (!key.isValid()) { continue; }

                // Handle the event
                if (key.isAcceptable()) { handleAccept(key); }
                } else if (key.isConnectable()) { handleConnect(key); }
                } else if (key.isReadable()) { handleRead(key); }
                } else if (key.isWritable()) { handleWrite(key); }
            } } catch (Exception e) { ... }
    }
}
```

Fabienne Boyer, Basics of Distributed Programming

5

NIO – Connect Handling

```
private void handleConnect(SelectionKey key) throws IOException {
    SocketChannel socketChannel = (ServerSocketChannel) key.channel();

    // finish establishing the connection
    socketChannel.finishConnect();

    // register the read interest on the selector
    socketChannel.register(this.selector, SelectionKey.OP_READ);

    ..
}
```

Fabienne Boyer, Basics of Distributed Programming

6

NIO – Accept Handling

```
private void handleAccept(SelectionKey key) throws IOException {

    ServerSocketChannel serverSocketChannel = (ServerSocketChannel) key.channel();

    // Accept the connection and make it non-blocking
    SocketChannel socketChannel = serverSocketChannel.accept();
    socketChannel.configureBlocking(false);

    // Register the new SocketChannel with our Selector, indicating
    // we'd like to be notified when there's data waiting to be read
    socketChannel.register(this.selector, SelectionKey.OP_READ);
}
```

Fabienne Boyer, Basics of Distributed Programming

7

NIO - Read Handling (basic version)

```
private void handleRead(SelectionKey key) throws IOException {
    SocketChannel socketChannel = (SocketChannel) key.channel();
    ByteBuffer inBuffer = ByteBuffer.allocate(length); // Read up to length bytes

    int nread = 0;
    try {
        nread = socketChannel.read(inBuffer);
    } catch (IOException e) {
        // the connection as been closed unexpectedly, cancel the selection and close the channel
        key.cancel();
        socketChannel.close();
        return;
    }
    if (nread == -1) {
        // the socket has been shutdown remotely cleanly
        key.channel().close();
        key.cancel();
        return;
    }
    // process the received data, being aware that it may be incomplete
    deliver(this, socketChannel, inBuffer.array(), nread);
}
```

Fabienne Boyer, Basics of Distributed Programming

8

NIO – Writing (basic version)

```
// outBuffers contains the data to write per channel
Hashtable<SocketChannel, ByteBuffer> outBuffers = ...;

private void write(SocketChannel socketChannel, byte[] data) throws IOException {

    // we suppose that previous data in outBuffer have already been sent
    // or we do not mind losing them
    outBuffers.put(socketChannel, ByteBuffer.wrap(data));

    // indicate we want to select OP-WRITE from now
    SelectionKey key = socketChannel.keyFor(this.selector);
    key.interestOps(SelectionKey.OP_READ | SelectionKey.OP_WRITE);
}
```

Fabienne Boyer, Basics of Distributed Programming

9

NIO – Basic Write Handling

```
// outBuffers contains the data to write per channel
Hashtable<SocketChannel, ByteBuffer> outBuffers = ...;

private void handleWrite(SelectionKey key) throws IOException {
    SocketChannel socketChannel = (SocketChannel) key.channel();
    ByteBuffer outBuffer = outBuffers.get(socketChannel);
    if (outBuffer.remaining() > 0) {
        try {
            socketChannel.write(outBuffer);
        } catch (IOException e) {
            // The channel has been closed
            key.cancel();
            socketChannel.close();
            return;
        }
    } else ...
}
```

Fabienne Boyer, Basics of Distributed Programming

10

NIO Buffers

- Attributes
 - Position (next index for read/write)
 - Limit (maximum number of bytes that can be read / written)
- Useful methods
 - **remaining()** always returns the value of limit - position
 - `ByteBuffer buf = ByteBuffer.wrap (byte[] b)`
assigns b as the buffer content, set (write) position to 0 and limit to b.length
 - `socketChannel.write(buf)`
send what can be send, update (write) position
 - `ByteBuffer buf = ByteBuffer.allocate(128);`
set (read) position to 0 and limit to 128
 - `socketChannel.read(buf)`
read what can be read, update (read) position
 - `Byte[] b = buf.array()`
assigns b with the buffer content
 - `buf.clear()`
set (read / write) position to 0

Fabienne Boyer, Basics of Distributed Programming

Advanced Reading & Writing

- For any channel, we manage a read automata
`Hashtable<SocketChannel, ReadAutomata> readAutomata = ...;`
- Read automata
 - Implements `handleRead(..)` to gather received bytes as they become available
 - Knows that each message is prefixed with its length on 4 bytes
 - Only deliver complete messages
- For any channel, we manage a write automata
`Hashtable<SocketChannel, WriteAutomata> writeAutomata =`
- Write automata
 - Manage a FIFO queue of messages to send
 - Prefix each message with its length when sending it
 - So each buffer to send contains the length of the message on 4 bytes

Fabienne Boyer, Basics of Distributed Programming